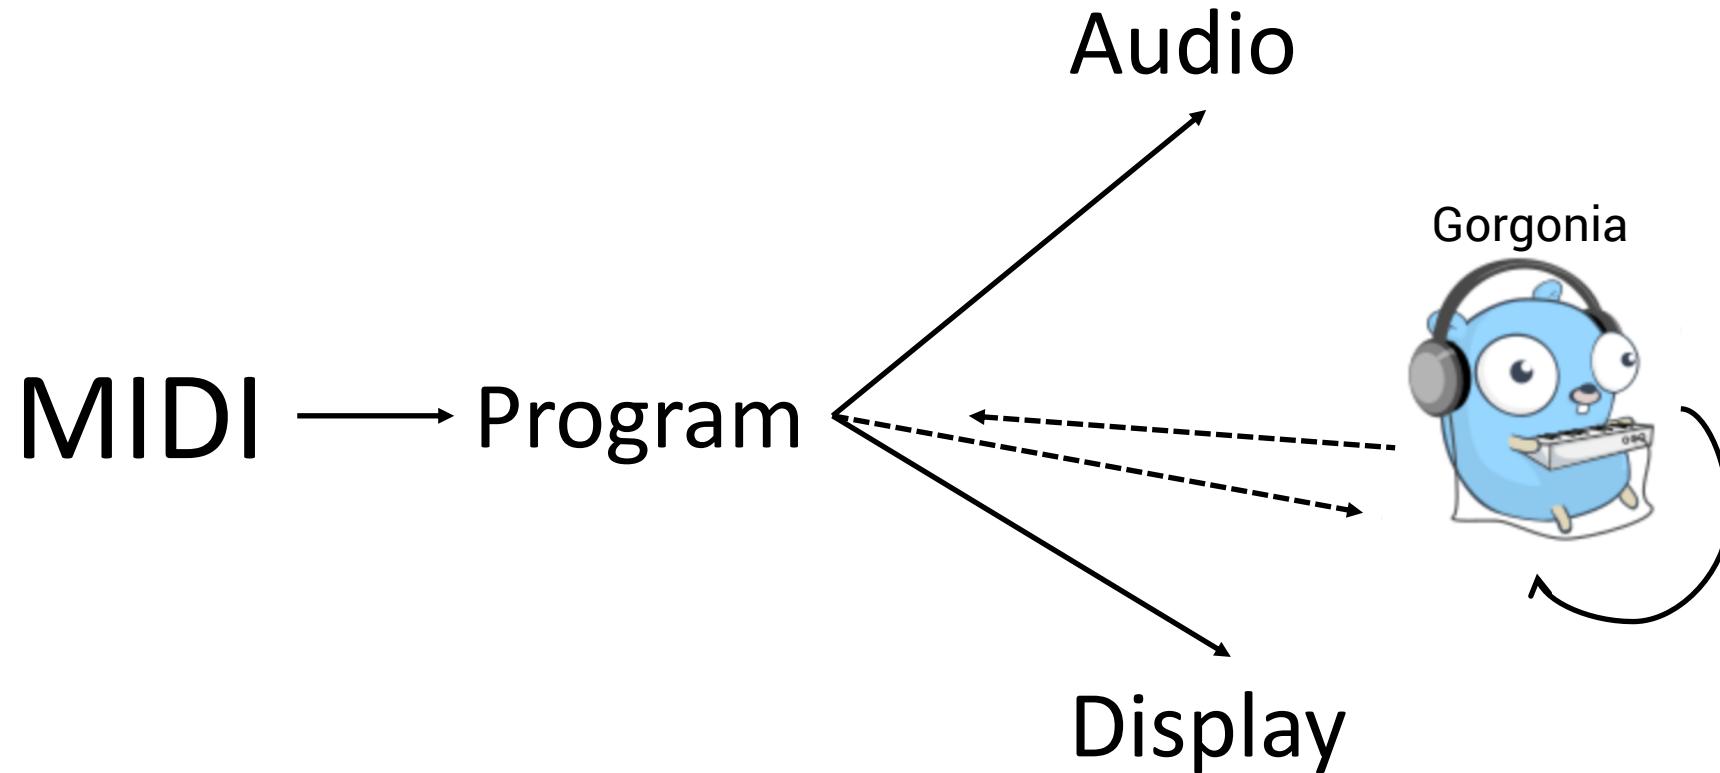

THE ART OF BONDAGE

Xuanyi Chew

GopherCon Singapore 2018

THE PROGRAM



[Download cuDNN v7.1.3 \[April 17, 2018\], for CUDA 9.1](#)

[cuDNN v7.1.3 Library for Linux](#)

[cuDNN v7.1.3 Library for Linux \(Power8/Power9\)](#)

[cuDNN v7.1.3 Library for Windows 7](#)

[cuDNN v7.1.3 Library for Windows 10](#)

[cuDNN v7.1.3 Runtime Library for Ubuntu16.04 \(Deb\)](#)

[cuDNN v7.1.3 Developer Library for Ubuntu16.04 \(Deb\)](#)

[cuDNN v7.1.3 Code Samples and User Guide for Ubuntu16.04 \(Deb\)](#)

[cuDNN v7.1.3 Runtime Library for Ubuntu16.04 & Power8 \(Deb\)](#)

[cuDNN v7.1.3 Developer Library for Ubuntu16.04 & Power8 \(Deb\)](#)

[cuDNN v7.1.3 Code Samples and User Guide for Ubuntu16.04 & Power8 \(Deb\)](#)

[cuDNN v7.1.3 Runtime Library for Ubuntu14.04 \(Deb\)](#)

[cuDNN v7.1.3 Developer Library for Ubuntu14.04 \(Deb\)](#)

[cuDNN v7.1.3 Code Samples and User Guide for Ubuntu14.04 \(Deb\)](#)

MACS ARE FRIENDS TOO!

:(
:(

WHY?

Follow @chewxy on Twitter

WHY?

- Existing libraries written in C/Fortran

WHY?

- Existing libraries written in C/Fortran
 - Difficult to port over to Go

WHY?

- Existing libraries written in C/Fortran
 - Difficult to port over to Go
- Our motivating example: CUDA libraries
 - import "gorgonia.org/cu"
 - import "gorgonia.org/cu/blas"
 - import "gorgonia.org/cu/dnn"

SOME DEFINITIONS

C API

```
cudnnStatus_t cudnnCreate(cudnnHandle_t *handle)
```

SOME DEFINITIONS

C API

```
cudnnStatus_t cudnnCreate(cudnnHandle_t *handle)
```

cgo Call

```
C.cudnnCreate(h *C.cudnnHandle_t) C.cudnnStatus_t
```

SOME DEFINITIONS

C API

```
cudnnStatus_t cudnnCreate(cudnnHandle_t *handle)
```

cgo Call

```
C.cudnnCreate(h *C.cudnnHandle_t) C.cudnnStatus_t
```

Go API

```
func New(h *C.cudnnHandle_t) C.cudnnStatus_t
func New(h *C.cudnnHandle_t) error
func New() (*Handle, error)
func New() *Handle
```

SOME DEFINITIONS

C API

```
cudnnStatus_t cudnnCreate(cudnnHandle_t *handle)
```

cgo Call

```
C.cudnnCreate(h *C.cudnnHandle_t) C.cudnnStatus_t
```

Go API

```
func New(h *C.cudnnHandle_t) C.cudnnStatus_t
func New(h *C.cudnnHandle_t) error
func New() (*Handle, error)
func New() *Handle
```

WHY WRITE AN API LIBRARY?

Follow @chewxy on Twitter

WHY WRITE AN API LIBRARY?

- A library that exposes a Go API reduces cognitive overhead.

WHY WRITE AN API LIBRARY?

- A library that exposes a Go API reduces cognitive overhead.
 - You don't have to know C anymore.

WHY WRITE AN API LIBRARY?

- A library that exposes a Go API reduces cognitive overhead.
 - You don't have to know C anymore.
- Abstract away bookkeeping activities.

WHY WRITE AN API LIBRARY?

- A library that exposes a Go API reduces cognitive overhead.
 - You don't have to know C anymore.
- Abstract away bookkeeping activities.
- More user friendly.

WHY WRITE AN API LIBRARY?

- A library that exposes a Go API reduces cognitive overhead.
 - You don't have to know C anymore.
- Abstract away bookkeeping activities.
- More user friendly.
- **Provide performance improvements for common use cases.**

WHY WRITE AN API LIBRARY?

In short, provide a great user experience.

THIS TALK

- Partly about strategies to write bindings.
- Partly about API design.

THE PROCESS: MANUAL WRITING

Follow @chewxy on Twitter

THE PROCESS: TRANSFORM FROM A TO B

cgo Call

```
C.cudnnCreate(h *C.cudnnHandle_t) C.cudnnStatus_t
```

THE PROCESS: TRANSFORM FROM A TO B

cgo Call

```
C.cudnnCreate(h *C.cudnnHandle_t) C.cudnnStatus_t
```



Wrap in one of these

Go API

```
func New(h *C.cudnnHandle_t) C.cudnnStatus_t  
func New(h *C.cudnnHandle_t) error  
func New() (*Handle, error)  
func New() *Handle
```

SIMPLEST EXAMPLE

```
cudnnStatus_t cudnnCreate(cudnnHandle_t *handle)
```



```
func New(handle: inout C.cudnnHandle_t) C.cudnnStatus_t
```

SIMPLEST EXAMPLE

```
func New(handle *C.cudnnHandle_t) C.cudnnStatus_t {  
    return C.cudnnCreate(handle)  
}
```

SIMPLEST EXAMPLE

```
func New(handle *C.cudnnHandle_t) C.cudnnStatus_t {  
    return C.cudnnCreateHandle()  
}
```

Terrible
Horrible
Vegetable

THE PROCESS

```
type Handle struct{
    internal C.cudnnHandle_t
}

func New() *Handle {
    var h C.cudnnHandle_t
    if st := cudnnCreate(&h); st != C.CUDNN_STATUS_SUCCESS {
        panic(st)
    }
    return &Handle{h}
}
```

THE PROCESS

```
type Handle C.cudnnHandle_t

func New() Handle {
    var h C.cudnnHandle_t
    if st := cudnnCreate(&h); st != C.CUDNN_STATUS_SUCCESS {
        panic(st)
    }
    return Handle(h)
}
```

THE PROCESS: GO HAS METHODS

- Methods make for nice APIs.

THE PROCESS: Go has METHODS

- Methods make for nice APIs.
- Methods allow for interfaces to be implemented.

MANUAL WRITING

- Forces one to actually think about the API.

MANUAL WRITING

- Forces one to actually think about the API.
 - Result: nicer API based on use case

MANUAL WRITING

- Forces one to actually think about the API.
- Requires one to be deeply familiar with the library.

GOOD EXAMPLES

- github.com/rakyll/portmidi
- github.com/mattn/go-gtk

MANUAL WRITING: DOWNSIDES

- A lot of time and effort required.

MANUAL WRITING: DOWNSIDES

- A lot of time and effort required.
- Not consistently replicable.

MANUAL WRITING: DOWNSIDES

- A lot of time and effort required.
- Not consistently replicable.
- Coverage issues.

MANUAL WRITING: DOWNSIDES

- A lot of time and effort required.
- Not consistently replicable.
- Coverage issues.
- Inability to keep up with new versions.



Wait a minute...
We are programmers!

Follow @chewxy on Twitter

AUTOMATIC GENERATION

Follow @chewxy on Twitter

AUTOMATIC GENERATION

- Repeatable results.

AUTOMATIC GENERATION

- Repeatable results.
- Generate multiple API levels.

AUTOMATIC GENERATION

Automatic generation steps:

1. Parse the C-header.
2. Generate the Go source code.

GOOD EXAMPLES

- github.com/go-gl/glfw
- github.com/go-gl/gl/vX.X-core/gl
 - Uses GLOW: github.com/go-gl/glow
- github.com/martine/gocairo
 - Uses rsc.io/c2go

IMMEDIATELY OBVIOUS SOLUTION

github.com/xlab/c-for-go

THE DEVIL IS IN THE DETAILS

- Enums
- Types
- Functions
- Procedures

THE DEVIL IS IN THE DETAILS

- Enums → Type + Const
- Types
- Functions
- Procedures

THE DEVIL IS IN THE DETAILS

- Enums → Type + Const
- Types → Type + Getter functions
- Functions
- Procedures

THE DEVIL IS IN THE DETAILS

- Enums → Type + Const
- Types → Type + Getter functions
- Functions → Functions/Methods
- Procedures

THE DEVIL IS IN THE DETAILS

- Enums → Type + Const
- Types → Type + Getter functions
- Functions → Functions/Methods.
- Procedures → Functions/Methods

NOTE FROM XLAB

"The resulting bindings are as low-level as C code, i.e. it would require knowledge of memory management to carefully use the resulting code, however no more C code is needed to make things done. ... Also usually a high-level wrapper is created by hand"

SOLUTION: MANUAL ANNOTATION

Annotate the things from the header – what should be a function, what should be a method, etc.

A BALANCE

Time spent manually annotating functions vs time spent generating code.

SOLUTION: MULTI-STAGE GENERATION

1. Generate annotations that can be deduced by rules.
2. Manually check/edit select annotations.
3. Generate library.

CUDA LIBRARIES

- Generated with:
 - github.com/cznic/cc
 - github.com/gorgonia/bindgen
 - [go/token](https://github.com/gorgonia/go-token)
 - [go/parser](https://github.com/gorgonia/go-parser)
 - [text/template](https://github.com/gorgonia/text-template)
 - Specific programs to generate each library:
 - gorgonia.org/cu/cmd/genlib
 - gorgonia.org/cu/cmd/gencublas
 - gorgonia.org/cu/cmd/gencudnn

TACTICAL ISSUES WITH AUTOMATIC GENERATION

Follow @chewxy on Twitter

ISSUES WITH AUTOMATIC GENERATION

- Unclear type

UNCLEAR TYPES

What is the type of labels?

```
cudnnStatus_t cudnnGetCTCLossWorkspaceSize(  
    cudnnHandle_t                      handle,  
    const cudnnTensorDescriptor_t        probsDesc,  
    const cudnnTensorDescriptor_t        gradientsDesc,  
    const int                           *labels,  
    const int                           *labelLengths,  
    const int                           *inputLengths,  
    cudnnCTCLossAlgo_t                 algo,  
    const cudnnCTCLossDescriptor_t     ctcLossDesc,  
    size_t                            *sizeInBytes)
```

UNCLEAR TYPES

What is the type of labels?

```
func (ctx *Handle) CTCLossWorkspaceSize(probDesc, gradientsDesc Tensor,  
                                         labels ???,  
                                         labelLengths ???,  
                                         inputLengths ???,  
                                         algo CTCLossAlgo,  
                                         ctcLossDesc Tensor) (size uintptr)
```

UNCLEAR TYPES

What is the type of labels?

```
cudnnStatus_t cudnnGetCTCLossWorkspaceSize(  
    cudnnHandle_t                      handle,  
    const cudnnTensorDescriptor_t        probsDesc,  
    const cudnnTensorDescriptor_t        gradientsDesc,  
    const int                           *labels,  
    const int                           *labelLengths,  
    const int                           *inputLengths,  
    cudnnCTCLossAlgo_t                 algo,  
    const cudnnCTCLossDescriptor_t     ctcLossDesc,  
    size_t                            *sizeInBytes)
```

UNCLEAR TYPES

Compare: What is the type of u , v?

```
cudnnStatus_t cudnnGetConvolution2dDescriptor(  
    const cudnnConvolutionDescriptor_t convDesc,  
    int                                *pad_h,  
    int                                *pad_w,  
    int                                *u,  
    int                                *v,  
    int                                *dilation_h,  
    int                                *dilation_w,  
    cudnnConvolutionMode_t             *mode,  
    cudnnDataType_t                   *computeType)
```

UNCLEAR TYPES

What is the type of labels?

```
cudnnStatus_t cudnnGetCTCLossWorkspaceSize(  
    cudnnHandle_t                      handle,  
    const cudnnTensorDescriptor_t        probsDesc,  
    const cudnnTensorDescriptor_t        gradientsDesc,  
    const int                           *labels,  
    const int                           *labelLengths,  
    const int                           *inputLengths,  
    cudnnCTCLossAlgo_t                 algo,  
    const cudnnCTCLossDescriptor_t     ctcLossDesc,  
    size_t                            *sizeInBytes)
```

UNCLEAR TYPES

What is the type of labels?

```
cudnnStatus_t cudnnGetCTCLossWorkspaceSize(  
    cudnnHandle_t                               handle,  
    const cudnnTensorDescriptor_t               probsDesc,  
    const cudnnTensorDescriptor_t               gradientsDesc,  
    const size_t                                labelsSize,  
    const int                                    labels[],  
    const int                                    labelLengths[],  
    const int                                    inputLengths[],  
    cudnnCTCLossAlgo_t                          algo,  
    const cudnnCTCLossDescriptor_t             ctcLossDesc,  
    size_t                                     *sizeInBytes)
```

UNCLEAR TYPES

What is the type of labels?

```
cudnnStatus_t cudnnGetCTCLossWorkspaceSize(  
    cudnnHandle_t                               handle,  
    const cudnnTensorDescriptor_t               probsDesc,  
    const cudnnTensorDescriptor_t               gradientsDesc,  
    const size_t                                labelsSize,  
    const int                                    labels[],  
    const int                                    labelLengths[],  
    const int                                    inputLengths[],  
    cudnnCTCLossAlgo_t                          algo,  
    const cudnnCTCLossDescriptor_t             ctcLossDesc,  
    size_t                                     *sizeInBytes)
```

To a C compiler, there is no difference between `int*` and `int[]`

ISSUES WITH AUTOMATIC GENERATION

- Unclear type
- Unclear use cases

UNCLEAR USE CASE

What is the use case of algo?

```
cudnnStatus_t cudnnGetCTCLossWorkspaceSize(  
    cudnnHandle_t                      handle,  
    const cudnnTensorDescriptor_t        probsDesc,  
    const cudnnTensorDescriptor_t        gradientsDesc,  
    const int                           *labels,  
    const int                           *labelLengths,  
    const int                           *inputLengths,  
    cudnnCTCLossAlgo_t                 algo,  
    const cudnnCTCLossDescriptor_t     ctcLossDesc,  
    size_t                            *sizeInBytes)
```

UNCLEAR USE CASE

What is the use case of algo?

```
cudnnStatus_t cudnnGetCTCLossWorkspaceSize(  
    cudnnHandle_t                      handle,  
    const cudnnTensorDescriptor_t        probsDesc,  
    const cudnnTensorDescriptor_t        gradientsDesc,  
    const int                           *labels,  
    const int                           *labelLengths,  
    const int                           *inputLengths,  
    const cudnnCTCLossAlgo_t           algo,  
    const cudnnCTCLossDescriptor_t      ctcLossDesc,  
    size_t                            *sizeInBytes)
```

HYBRID SOLUTION

1. Parse the docs. Add them as constants.
2. Manually annotate specific parameters.
3. Automatic generation as an aid to manual generation.

VARIOUS CONSIDERATIONS

Follow @chewxy on Twitter

VARIOUS CONSIDERATIONS

- Performance
- Contexts and multi-level API
- Error detection
- Reporting

VARIOUS CONSIDERATIONS

- Performance
- Contexts and multi-level API
- Error detection
- Reporting

PERFORMANCE TIP: BATCHING

```
func NewConvolution(mathType MathType, groupCount int, padding, filterStride, dilation []int, convolutionMode ConvolutionMode, datatype DataType) (*Convolution, error) {
    var internal C.cudnnConvolutionDescriptor_t
    if err := result(C.cudnnCreateConvolutionDescriptor(&internal)); err != nil { ... }
    if err := result(C.cudnnSetConvolutionMathType(internal, mathType.C())); err != nil { ... }
    if err := result(C.cudnnSetConvolutionGroupCount(internal, C.int(groupCount))); err != nil { ... }
    switch len(padding) {
        case 0, 1:
            return nil, errors.Errorf("Only 2+ dims are allowed")
        case 2:
            // snipped: destructuring of padH, padW, u, v, etc
            if err := result(C.cudnnSetConvolution2dDescriptor(internal, C.int(padH), C.int(padW), C.int(u), C.int(v), C.int(dilationH),
C.int(dilationW), convolutionMode.C(), datatype.C())); err != nil {
                return nil, err
            }
        default:
            // snipped: converting padding, filterStride, dilation to a pointer
            if err := result(C.cudnnSetConvolutionNdDescriptor(internal, C.int(len(padding)), padA, strideA, dilationA,
convolutionMode.C(), datatype.C())); err != nil {
                return nil, err
            }
    }
    retVal := &Convolution{...}
    runtime.SetFinalizer(retVal, destroyConvolution)
    return retVal, nil
}
```

PERFORMANCE TIP: BATCHING

```
func NewConvolution(mathType MathType, groupCount int, padding, filterStride, dilation []int, convolutionMode ConvolutionMode, datatype  
    dataType) (*Convolution, error) {  
    var internal C.cudnnConvolutionDescriptor_t  
    if err := result(C.cudnnCreateConvolutionDescriptor(&internal)); err != nil { ... }  
    if err := result(C.cudnnSetConvolutionMathType(internal, mathType.C())); err != nil { ... }  
    if err := result(C.cudnnSetConvolutionGroupCount(internal, C.int(groupCount))); err != nil { ... }  
    switch len(padding) {  
        case 0, 1:  
            return nil, errors.Errorf("Only 2+ dims are allowed")  
        case 2:  
            // snipped: destructuring of padH, padW, u, v, etc  
            if err := result(C.cudnnSetConvolution2dDescriptor(internal, C.int(padH), C.int(padW), C.int(u), C.int(v), C.int(dilationH),  
                C.int(dilationW), convolutionMode.C(), datatype.C())); err != nil {  
                return nil, err  
            }  
        default:  
            // snipped: converting padding, filterStride, dilation to a pointer  
            if err := result(C.cudnnSetConvolutionNdDescriptor(internal, C.int(len(padding)), padA, strideA, dilationA,  
                convolutionMode.C(), datatype.C())); err != nil {  
                return nil, err  
            }  
    }  
    retVal := &Convolution{...}  
    runtime.SetFinalizer(retVal, destroyConvolution)  
    return retVal, nil  
}
```

PERFORMANCE TIP: BATCHING

```
cudnnStatus_t gocudnnNewConvolution(cudnnConvolutionDescriptor_t *retVal, cudnnMathType_t mathType, const int groupCount,
                                     const int size, const int* padding, const int* filterStrides, const int* dilation, cudnnConvolutionMode_t convolutionMode,
                                     cudnnDataType_t dataType) {
    cudnnStatus_t status ;
    status = cudnnCreateConvolutionDescriptor(retVal);
    if (status != CUDNN_STATUS_SUCCESS) { return status; }

    status = cudnnSetConvolutionMathType(*retVal, mathType);
    if (status != CUDNN_STATUS_SUCCESS) { return status; }

    status = cudnnSetConvolutionGroupCount(*retVal, groupCount);
    if (status != CUDNN_STATUS_SUCCESS) { return status; }

    switch (size) {
        case 0, 1 :
            return CUDNN_STATUS_BAD_PARAM;
        case 2:
            // snipped : unpacking the params
            status = cudnnSetConvolution2dDescriptor(*retVal, padH, padW, u, v, dilationH, dilationW, convolutionMode, dataType);
            break;
        default:
            status = cudnnSetConvolutionNdDescriptor(*retVal, paddingSize, padding, filterStrides, dilation, convolutionMode,
dataType);
            break;
    }
    return status;
}
```

PERFORMANCE TIP: BATCHING

```
func NewConvolution(mathType MathType, groupCount int, padding, filterStride, dilation []int, convolutionMode ConvolutionMode, datatype  
    dataType) (*Convolution, error) {  
    var internal C.cudnnConvolutionDescriptor_t  
    if err := result(C.cudnnCreateConvolutionDescriptor(&internal)); err != nil { ... }  
    if err := result(C.cudnnSetConvolutionMathType(internal, mathType.C())); err != nil { ... }  
    if err := result(C.cudnnSetConvolutionGroupCount(internal, C.int(groupCount))); err != nil { ... }  
    switch len(padding) {  
        case 0, 1:  
            return nil, errors.Errorf("Only 2+ dims are allowed")  
        case 2:  
            // snipped: destructuring of padH, padW, u, v, etc  
            if err := result(C.cudnnSetConvolution2dDescriptor(internal, C.int(padH), C.int(padW), C.int(u), C.int(v), C.int(dilationH),  
                C.int(dilationW), convolutionMode.C(), datatype.C())); err != nil {  
                    return nil, err  
                }  
        default:  
            // snipped: converting padding, filterStride, dilation to a pointer  
            if err := result(C.cudnnSetConvolutionNdDescriptor(internal, C.int(len(padding)), padA, strideA, dilationA,  
                convolutionMode.C(), datatype.C())); err != nil {  
                    return nil, err  
                }  
    }  
    retVal := &Convolution{...}  
    runtime.SetFinalizer(retVal, destroyConvolution)  
    return retVal, nil  
}
```

PERFORMANCE TIP: BATCHING

```
func NewConvolution(mathType MathType, groupCount int, padding, filterStride, dilation []int, convolutionMode ConvolutionMode, datatype DataType) (*Convolution, error) {
    // elided: checks

    var internal C.cudnnConvolutionDescriptor_t
    padA := (*C.int)(unsafe.Pointer(&padding[0]))
    strideA := (*C.int)(unsafe.Pointer(&filterStride[0]))
    dilationA := (*C.int)(unsafe.Pointer(&dilation[0]))
    if err = result(C.gocudnnNewConvolution(&internal, mathType.C(), C.int(groupCount), C.int(len(padding)), padA, filterStrideA,
dilationA, convolutionMode.C(), datatype.C())); err != nil {
        return nil, err
    }

    retVal := &Convolution{...}
    runtime.SetFinalizer(retVal, destroyConvolution)
    return retVal, nil
}
```

PERFORMANCE: BATCHING TIPS

- astyle is your friend.
- astyle --style=google \
 --lineend=linux \
 --indent=tab \
 --indent-switches \
 --align-pointer=type \
 --align-reference=name \
 --delete-empty-lines
- <http://astyle.sourceforge.net/>

PERFORMANCE TIP: CACHE FIELDS

Opaque types:

```
struct cudnnTensorStruct* cudnnTensorDescriptor_t
```

Go version:

```
type TensorDescriptor struct {
    internal C.cudnnTensorDescriptor_t

    format TensorFormat
    dataType DataType
    shape []int
    strides []int
}
```

PERFORMANCE TIP: CACHE FIELDS

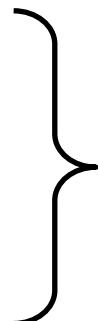
Opaque types:

```
struct cudnnTensorStruct* cudnnTensorDescriptor_t
```

Go version:

```
type TensorDescriptor struct {
    internal C.cudnnTensorDescriptor_t

    format TensorFormat
    dataType DataType
    shape []int
    strides []int
}
```



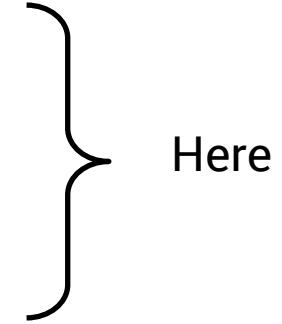
Where did these come from?

PERFORMANCE TIP: CACHE FIELDS

```
cudnnStatus_t cudnnSetTensor4dDescriptor(  
    cudnnTensorDescriptor_t          tensorDesc,  
    cudnnTensorFormat_t              format,  
    cudnnDataType_t                 dataType,  
    int                             n,  
    int                             c,  
    int                             h,  
    int) ;
```

PERFORMANCE TIP: CACHE FIELDS

```
cudnnStatus_t cudnnSetTensor4dDescriptor(  
    cudnnTensorDescriptor_t          tensorDesc,  
    cudnnTensorFormat_t              format,  
    cudnnDataType_t                 dataType,  
    int                             n,  
    int                             c,  
    int                             h,  
    int) ;
```



Here

PERFORMANCE TIP: CACHE FIELDS

```
cudnnStatus_t cudnnSetTensor4dDescriptor(  
    cudnnTensorDescriptor_t tensorDesc,  
    cudnnTensorFormat_t format,  
    cudnnDataType_t dataType,  
    int n,  
    int C,  
    int h,  
    int w);  
  
cudnnStatus_t cudnnGetTensor4dDescriptor(  
    const cudnnTensorDescriptor_t tensorDesc,  
    cudnnDataType_t *dataType,  
    int *n,  
    int *C,  
    int *h,  
    int *w,  
    int *nStride,  
    int *cStride,  
    int *hStride,  
    int *wStride );
```

PERFORMANCE TIP: CACHE FIELDS

```
cudnnStatus_t cudnnSetTensor4dDescriptor(  
    cudnnTensorDescriptor_t tensorDesc,  
    cudnnTensorFormat_t format,  
    cudnnDataType_t dataType,  
    int n,  
    int C,  
    int h,  
    int w);  
  
cudnnStatus_t cudnnGetTensor4dDescriptor(  
    const cudnnTensorDescriptor_t tensorDesc,  
    cudnnDataType_t *dataType,  
    int *n,  
    int *C,  
    int *h,  
    int *w,  
    int *nStride,  
    int *cStride,  
    int *hStride,  
    int *wStride );
```

GUIDING PRINCIPLE

- Make your data structures useful.

GUIDING PRINCIPLE

- Make your data structures useful.
- Start from the use cases.

GUIDING PRINCIPLE

- Make your data structures useful.
- Start from the use cases.
- Eliminate binding calls where unnecessary.

GUIDING PRINCIPLE

- Make your data structures useful.
- Start from the use cases.
- Eliminate binding calls where unnecessary.
- Minimize errors – check them in Go before passing off to cgo.

VARIOUS CONSIDERATIONS

- Performance
- Contexts and multi-level API
- Error detection
- Reporting

CONTEXTUAL USE

Most graphics heterogenous memory related libraries require a thread-lock.

CONTEXTUAL USE: GENERATE MULTI-LEVEL API

- One that handles the common use case.
- One low level API, leaving it to the user to manage themselves.

CUDA CONTEXTS

```
// CUContext is a CUDA context
type CUContext uintptr
```

CUDA CONTEXTS

```
// CUContext is a CUDA context
type CUContext uintptr

// Ctx is a standalone CUDA Context that is threadlocked.
type Ctx struct{
    CUContext
    // contains filtered or unexported fields
}
```

CUDA CONTEXTS

```
// CUContext is a CUDA context
type CUContext uintptr

// Ctx is a standalone CUDA Context that is threadlocked.
type Ctx struct{
    CUContext
    // contains filtered or unexported fields
}

// BatchedContext is a CUDA context where the CUDA calls are batched up.
type BatchedContext struct {
    Context
    Device
    // contains filtered or unexported fields
}
```

CUDA CONTEXT: SYNCHRONIZE()

```
func Synchronize() (err error) {  
    return result(C.cuCtxSynchronize())  
}
```

Use in a main-threadlocked function

```
...  
handleErr(ctx.Lock())  
handleErr(cu.SetCurrentContext(ctx))  
handleErr(cu.Synchronize())  
handleErr(ctx.Unlock())
```

CUDA CONTEXT: SYNCHRONIZE()

```
// Ctx is a standalone CUDA Context that is threadlocked.  
type Ctx struct {  
    work    chan (func() error)  
    ...  
}  
  
func (ctx *Ctx) Do(fn func() error) error {  
    ctx.work <- fn  
    return <-ctx.errChan  
}  
  
func (ctx *Ctx) Synchronize() {  
    f := func() error {  
        return result(C.cuCtxSynchronize())  
    }  
    ctx.err = ctx.Do(f)  
}
```

VARIOUS CONSIDERATIONS

- Performance
- Contexts and multi-level API
- Error detection
- Reporting

ERROR DETECTION

- Leverage existing Go source code analysis tools:
 - goimports
 - govet
 - go/token, go/parser, go/ast

ERROR DETECTION

- Potential Nil-checks
 - Alias analysis is undecidable^[0].
 - Doesn't mean you can't try.
 - Easy to check when functions are small.

[0] G. Ramalingam (1994): The Undecidability of Aliasing.

VARIOUS CONSIDERATIONS

- Performance
- Contexts and multi-level API
- Documentation
- Error detection
- Reporting

REPORTING

- In multistage generation programs – always report back what is not done yet.

REPORTING

- In multistage generation programs – always report back what is not done yet.

<input type="checkbox"/>	! API for JIT modules	API Completion	help wanted
	#13 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuModuleGetSurfRef`	API Completion	help wanted
	#12 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuModuleGetTexRef`	API Completion	help wanted
	#11 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuModuleLoadFatBinary`	API Completion	help wanted
	#10 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuModuleLoadDataEx`	API Completion	help wanted
	#9 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuMemHostGetFlags`	API Completion	help wanted
	#8 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuMemHostGetDevicePointer`	API Completion	help wanted
	#7 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuMemHostAlloc`	API Completion	help wanted
	#6 opened on 14 May 2017 by chewxy		
<input type="checkbox"/>	! API for `cuMemAllocHost`	API Completion	help wanted
	#5 opened on 14 May 2017 by chewxy		

FINALLY...

Follow @chewxy on Twitter

P/s

The source for the demo is online:

<https://github.com/chewxy/gopherconsg2018>

IN THIS TALK...

- Introduced the *whys* of writing an API library for Go bindings.
- Showed *how* to think about writing an API library.
- Using gorgonia.org/cu/dnn and gorgonia.org/cu as a concrete examples.

THE ASK

Follow @chewxy on Twitter

THE Ask

- Through use and feedback, APIs can be improved.

THE Ask

- Through use and feedback, APIs can be improved.
- Please use these, and provide feedback:
 - gorgonia.org/gorgonia
 - gorgonia.org/tensor
 - [gorgonia.org/cu/...](https://gorgonia.org/cu/)

THE Ask

- Through use and feedback, APIs can be improved.
- Please use these, and provide feedback:
 - gorgonia.org/gorgonia
 - gorgonia.org/tensor
 - [gorgonia.org/cu/...](http://gorgonia.org/cu/)
- Help improve the generation programs:
 - gorgonia.org/cu/cmd/genlib
 - gorgonia.org/cu/cmd/gencudnn
 - gorgonia.org/cu/cmd/gencublas

Follow me on Twitter: @chewxy
Email me: chewxy@gmail.com

THANK YOU GOPHERCON SG!

Follow @chewxy on Twitter